

# 1. Pairwise Preference Generation Script

`build_pairwise_prefs.py`, generates **pairwise preference data** from a dataset of queries, code, and descriptions. These preference pairs (`prompt`, `chosen`, `rejected`) are useful for training **reward models** in RLHF (Reinforcement Learning with Human Feedback) pipelines or similar tasks.

---

## Overview

The script performs the following steps:

1. Retrieve relevant context for each query from a **vector database**.
  2. Generate multiple candidate responses using a **large language model (LLM)**.
  3. Score each candidate using a **grounding-first heuristic** combining context relevance, keyword coverage, and length penalty.
  4. Select the best and worst candidate as a pair: (`chosen`, `rejected`).
  5. Save the generated pairs in a JSONL file for downstream training.
- 

## Major Components

### 1. Configuration

- **DATA\_CSV**: Input CSV containing queries, code, and descriptions.
- **VDB\_PATH**: Path to the local FAISS vector store.
- **BASE\_LLM**: Pretrained model used for candidate generation.
- **USE\_4BIT**: Whether to use 4-bit quantization for memory efficiency.

- **MAX\_NEW\_TOKENS**: Maximum number of tokens generated per candidate.
  - **DEVICE**: "cuda" if GPU is available, otherwise "cpu".
  - **SEEDS**: Seeds for generating diverse candidates.
  - **NUM\_CANDIDATES**: Number of responses generated per prompt.
  - **OUT\_JSONL**: File to save the preference pairs.
- 

## 2. Retrieval

- Uses **FAISS** for efficient vector similarity search.
  - Embeddings are generated with a **sentence-transformer**.
  - Retrieves the top **k** documents (**k=1** in this case, since documents are dense/self-contained).
  - These retrieved documents form the **context** for the LLM.
- 

## 3. Prompt Construction

- The function `make_prompt(context, question)` formats the input for the LLM:
    - Includes retrieved context.
    - Adds the user query.
    - Wraps the text with instruction tags.
    - Ensures the LLM answers based on context or says "I don't know" if unsure.
- 

## 4. Model Loading

- Supports **4-bit quantized LLMs** for memory efficiency.
  - The Mistral-7B-Instruct model is loaded using `AutoModelForCausalLM`.
  - Optional **bfloat16** computations improve numeric stability.
  - Tokenizer is loaded and ensures a padding token is set for generation.
- 

## 5. Candidate Generation

- For each prompt, multiple candidates are generated for diversity.
  - Uses **sampling** with temperature and top-p (`do_sample=True, temperature=0.7, top_p=0.9`).
  - Multiple seeds ensure different outputs for the same prompt.
  - The function `extract_answer` extracts the generated text after the instruction tags.
  - Optimized batching reduces generation time significantly.
- 

## 6. Scoring Candidates

Each candidate is scored using a **heuristic function**:

1. **Grounding Score**: Cosine similarity between candidate and retrieved context.
  2. **Keyword Coverage**: Fraction of task-relevant keywords present in the candidate.
  3. **Length Penalty**: Mild penalty for overly long outputs.
- Scoring formula:  
$$\text{score} = (W_{\text{GROUND}} * \text{grounding}) + (W_{\text{KEYWD}} * \text{keyword\_coverage}) - (W_{\text{LEN}} * \text{length\_penalty})$$
  - Weights:  $W_{\text{GROUND}} = 0.7$ ,  $W_{\text{KEYWD}} = 0.3$ ,  $W_{\text{LEN}} = 0.05$ .

---

## 7. Main Loop: Building Preference Pairs

- Iterates through each query in the CSV.
- Retrieves context from FAISS.
- Generates multiple candidates.
- Scores candidates and sorts them.
- Selects **best (chosen)** and **worst (rejected)** candidate.
- Writes a JSON object per pair in the output JSONL file.

Output JSONL structure:

- **"prompt"**: LLM-formatted prompt
- **"chosen"**: Best candidate according to the heuristic
- **"rejected"**: Worst candidate according to the heuristic

---

## Key Features

- **Retrieval-Augmented Generation (RAG)**: Ground answers in relevant context.
  - **Diverse Candidate Generation**: Multiple seeds and sampling strategies.
  - **Heuristic Scoring**: Ensures contextually correct and relevant chosen answers.
  - **Memory Efficient**: 4-bit quantization and device-aware model loading.
  - **Extensible**: Adjust **NUM\_CANDIDATES**, seeds, or scoring weights to tune output.
-

## Notes

- Generating candidates for 500 queries takes roughly 3–4 hours on a GPU with 4-bit quantization.
- Tokenizer is ensured to have a padding token to avoid generation errors.
- Output JSONL can directly be used to train reward models for RLHF pipelines.